



# Lecture 7

Processes and threads.

Signals.

*Department of Computer Systems FIT, Czech Technical University in Prague*

*©Jan Trdlička, 2010*



- Process is **running program**.
- Each process is identified by **unique process identifier (PID)** used by kernel.
- Each process knows **its parent** under the **parent process identifier (PPID)**.
- New process is created by system call:
  - **fork()**
    - function creates a new process
    - the address space of the new process (child process) is an exact copy of the address space of the calling process (parent process)
    - function returns
      - -1 (error) or child PID in parent process
      - 0 in child process
    - child process has some new properties (PID, PPID,...) and other properties are inherited (e.g. EUID, EGID,...) from parent process
  - **exec()**
    - function replaces the current process image with a new process



# Example

```
int main (void)
{ ...
    pid = fork();
    switch (pid) {
    case -1: /* doslo k chybe */
        perror ("chyba ve funkci fork()");
        exit (1);
    case 0: /* kod provadeny v potomkovi */
        printf ("PID procesu potomka: %d\n", (int) getpid ());
        execlp("sleep", "sleep", "30", (char *) NULL);
        perror ("chyba ve funkci execlp()");
        exit (1);
    default: /* program provadeny v rodici */
        printf ("PID procesu rodice : %d\n", (int) getpid ());
        wait(&status);
    };
    ... }
```



- Thread is **running subprogram** in process/kernel environment.
- **Process containing  $n$  threads** can be executed **concurrently on  $n$  CPU's**.
- Process with one thread (default) can use only one CPU at a given time.
- **Thread creation is faster** then process creation.
- New thread can be created by library function, e.g. `pthread_create()`.



# Example

...

```
void *kod_vlakna(void *threadid)
{ printf("ID vlakna: %ld\n", (long int) threadid);
  sleep(60);
  pthread_exit(NULL);
}
```

...

```
int main(void)
{ pthread_t threads[NUM_THREADS];
  int rc, i;
  for(i=0; i<NUM_THREADS; i++){
    rc = pthread_create(&threads[i], NULL, kod_vlakna, (void *) i);
    if (rc){ perror("Chyba ve funkci pthread_create()"); exit(1); }
  }
  /* pthread_exit(NULL); */
  /* pthread_join(threads[i],&data); */
  ...}
```




# Context switching

- Computing process of storing and restoring state (context) of a CPU so that execution can be resumed from the same point at a later time
- This enables multiple threads/processes to share available CPUs.
- **Kernel determines when and who gets the CPU.**
- **Kernel uses several information to make decision**
  - Process/thread priority (e.g. 0 – 169 in Solaris)
  - Priority class (e.g. TS, IA, FSS, FX, SYS, RT in Solaris)
  - Thread state, thread behavior in history, ...
- **Thread/process priority** can be fixed or dynamic.
- Thread can use CPU only during some **time quantum**.
- **Size of time quantum** can be different for different threads/processes and can vary in time (depends on Unix implementation).




# Context switching


## Memory



## CPU 0 - usage

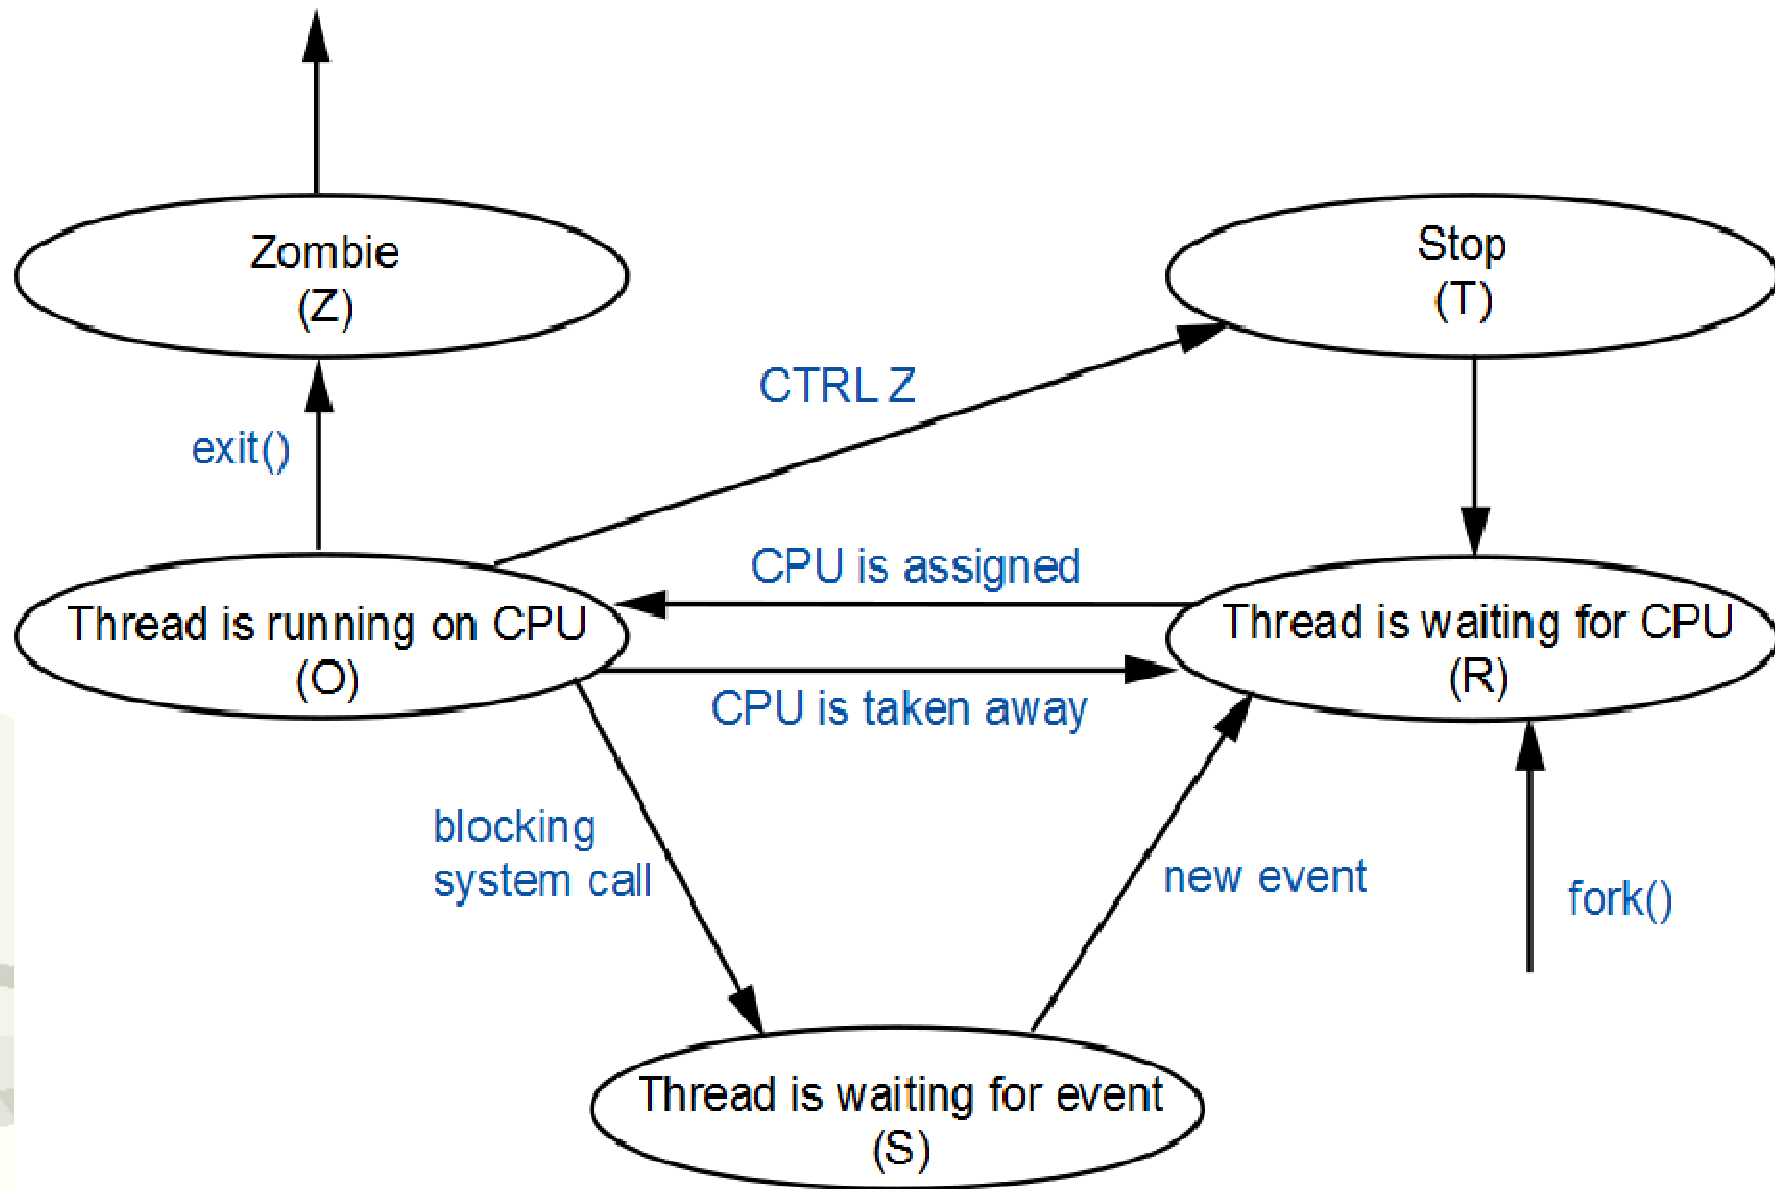


## CPU 1- usage





# Thread states





# Process/thread listing

`ps`

- prints information about processes that have the same effective user ID and the same controlling terminal as the invoker

`ps -e`

- Lists information about every running process

`ps -f` or `ps -l`

- Prints more details about running processes :

- S state(O, S, R, Z, T)
- PID, PPID process ID, parent process ID
- PRI priority
- NI NICE value
- STIME starting time
- TIME cumulative execution time
- TTY controlling terminal
- CMD command name

## `ps -o format`

- allows the output format to be specified under user control
- **format** specification must be a list of names:
  - user ruser group rgroup uid ruid gid rgid pid ppid pgid sid pri  
nice class time etime stime s c lwp ...

## `ps -Le`

- Prints information about each thread (light weight process)

`pgrep [-lvx] [pattern] [-u users...]`

- reports the process IDs of the processes whose attributes match the criteria specified on the command line
- `-l` Prints the process name along with the process ID of each matching process.
- `-v` Reverses the sense of the matching.

`prstat` or `top`

- iteratively examines all active processes on the system and reports statistics based on the selected output mode and sort order

`ptree [-a] [ pid ] [ user ]` in Solaris

`pstree [-a] [ pid ] [ user ]` in Linux

- prints the process trees containing the specified pids or users, with child processes indented from their respective parent processes



# Process priority

- **Process priority can be decreased (root can also increase) by command:**

`nice -priority program`

`nice -n priority program`

- where **priority** is integer number 1-19
- higher number = lower priority
- negative number = increasing of priority (only root)
- In Solaris: better command for priority modification is **prionctl**



# How to write a correct script?

- What is the difference in the following loops?

```
while [ -f $1 ] ; do : ; done; echo " $1 is removed "
```

```
while [ -f $1 ] ; do sleep 2 ; done; echo " $1 is removed "
```

- Which loop will determine earlier that the file is removed if the loop is executed concurrently 100x (every time with different file name \$1)?



- Limited form of inter-process communication used in Unix
- Kernel **interrupts the process's normal flow of execution**, when signal is sent to a process.
- **Signal can be send by**
  - kernel (e.g. arithmetic exception, segmentation fault...)
  - terminal driver (e.g. key sequence CTRL C, CTRL \ , ...)
  - jiným procesem (např. příkaz nebo funkce kill)
- Signal is identified by **name** and **number**.
- Reaction to signal: none, exit, exit+core
- **List of signals:**
  - Command `kill -l`
  - In Unix manual `man -s 3HEAD signal` (Solaris)
- **How to send signal:**

```
kill -signal PID
```

```
pkill -signal [ -vx] [ pattern -u users ...]
```



# Signals and terminal driver

- **Some signal can be sent by sequence of keys:**
- (see `stty -a` or `man stty`)

Key sequences	Meaning	Signal
CTRL C	Interrupt process	2 SIGINT
CTRL \	Quit process	3 SIGQUIT
CTRL Z	Suspend process (not in sh). <b>Not process termination!!!</b>	24 SIGTSTP





# Important signals

**15 SIGTERM (TERMinate)**

**9 SIGKILL (KILL)**

**2 SIGINT (INTerrupt)** and **3 SIGQUIT (QUIT)**

**1 SIGHUP (HangUP)**

- Parent process sends this signal to its child processes during termination
- If the child process must continue after parent process termination you must start children process by command

`nohup příkaz &`



# Reaction to signal

- **Default signal reaction** are set during process startup.
- Process can **modify signal reaction** by system call `trap( )` (except signals KILL and STOP).
- We can also modify signal reaction by command `trap`.
  - **Signal reaction setup**  
`trap 'commands' signals`
  - **Print definition of signal reaction:**  
`trap`
  - **Setup ignoring:**  
`trap ' ' signals`
  - **Setup default signal reaction:**  
`trap - signals`



# Command execution

- **In foreground**

**\$ command**

- Command (no built-in command) is executed like new process.
- Standard input and outputs are assigned to terminal..
- Shell waits for command termination.

- **In background**

**\$ command &**

- Command (no built-in command) is executed like new process.
- Only standard outputs are assigned to terminal (if the command try to read from stdin the it is stoped)
- Shell doesn't wait for command termination.



# Command execution

- **Immune to hang-ups**

`$ nohup command &`

- Command is running after shell termination.
- Standard outputs are redirected to file `nohup.out`.

- **In given time**

- Command can be started at a given time by commands `at` or `crontab`.
- System process `cron` executes the command under user identity at a given time.
- Standard outputs are sent by email.



- Properties of shell (except /bin/sh)
- **Every process which is running in given shell is assigned job identifier (JID) in this shell.**
- In the shell we can use the following commands
  - `jobs` list all running process inside this shell
  - `fg [%JID]` move process to foreground
  - `bg [%JID]` move process to foreground
  - `kill -signal [%JID]` send signal to process
- If JID is not specified, the last process is used.



# Useful commands

---

- **Elapsed time report**

`timex` command

- **System calls used by command**

`truss` command

